

# Berkstats Day 2

Ulrich Matter

22 February 2017

## Working with Data in R: Data Structures and Indices

### Vectors and Lists

```
# A vector containing numeric (or integer) values
numeric_vector <- 10:20
numeric_vector[2]
```

```
## [1] 11
```

```
numeric_vector[2:5]
```

```
## [1] 11 12 13 14
```

```
# A string vector ('a vector containing text')
string_vector <- c("a", "b", "c")
string_vector[-3]
```

```
## [1] "a" "b"
```

```
# Lists
```

```
# A list can contain different types of elements, for example a numeric vector and a string_vector
mylist <- list(numbers = numeric_vector, letters = string_vector)
mylist
```

```
## $numbers
```

```
## [1] 10 11 12 13 14 15 16 17 18 19 20
```

```
##
```

```
## $letters
```

```
## [1] "a" "b" "c"
```

```
# We can access the elements of a list in various ways
# with the element's name
```

```
mylist$numbers
```

```
## [1] 10 11 12 13 14 15 16 17 18 19 20
```

```
mylist["numbers"]
```

```
## $numbers
```

```
## [1] 10 11 12 13 14 15 16 17 18 19 20
```

```
# via the index
```

```
mylist[1]
```

```
## $numbers
```

```
## [1] 10 11 12 13 14 15 16 17 18 19 20
```

```
# with [[]] we can access directly the content of the element
```

```
mylist[[1]]
```

```
## [1] 10 11 12 13 14 15 16 17 18 19 20
```

```
# lists can also be nested (list of lists of lists....)
mynestedlist <- list(a = mylist, b = 1:5)
```

## Matrices and Data Frames

```
# matrices
mymatrix <- matrix(numeric_vector, nrow = 4)

## Warning in matrix(numeric_vector, nrow = 4): data length [11] is not a sub-
## multiple or multiple of the number of rows [4]

# get the second row
mymatrix[2,]

## [1] 11 15 19

# get the first two columns
mymatrix[, 1:2]

##      [,1] [,2]
## [1,]  10  14
## [2,]  11  15
## [3,]  12  16
## [4,]  13  17

# data frames ("lists as columns")
mydf <- data.frame(Name = c("Alice", "Betty", "Claire"), Age = c(20, 30, 45))
mydf

##      Name Age
## 1  Alice  20
## 2  Betty  30
## 3 Claire  45

# select the age column
mydf$Age

## [1] 20 30 45

mydf[, "Age"]

## [1] 20 30 45

mydf[, 2]

## [1] 20 30 45

# select the second row
mydf[2,]

##      Name Age
## 2  Betty  30
```

## Classes and Data Structure

```
# have a look at what kind of object you are dealing with
class(mydf)
```

```
## [1] "data.frame"
class(mymatrix)

## [1] "matrix"
# have a closer look at the data structure
str(mydf)


## 'data.frame':  3 obs. of  2 variables:
##  $ Name: Factor w/ 3 levels "Alice","Betty",...: 1 2 3
##  $ Age : num  20 30 45
```

## Z-Scores and the Standard Normal Distribution

```
# The Z-Score formula in R
# define the parameter values
X <- 10
mu <- 12
sigma <- 2
# compute the z-score
z <- (X - mu) / sigma
z

## [1] -1

# Plot the Standard Normal Distribution
plot(function(x) dnorm(x), -4, 4, main = "Normal density")
```



```
# Get the area under the curve (probability of observing a value of a certain size)  
pnorm(-1)
```

```
## [1] 0.1586553
```

```
pnorm(-2)
```

```
## [1] 0.02275013
```

```
pnorm(-1) - pnorm(-2)
```

```
## [1] 0.1359051
```

## Standard Errors

```
# formula for standard error  
# define parameter values  
s <- 20  
n <- 100  
# compute the standard error (of the mean)  
se <- s / sqrt(n)  
se
```

```
## [1] 2
```

```
# write your own standard error (of the mean) function  
se <- function(x) {  
  s <- sd(x)  
  n <- length(x)  
  se <- s / sqrt(n)  
  
  return(se)  
}
```

```
# draw a random sample of size 100 and compute the mean and its estimated standard error  
mysample <- rnorm(100)  
mean(mysample)
```

```
## [1] -0.09034929
```

```
se(mysample)
```

```
## [1] 0.09682044
```

```
# repeat this but this time with a larger sample  
mysample <- rnorm(1000)  
mean(mysample)
```

```
## [1] -0.01860827
```

```
se(mysample)
```

```
## [1] 0.03048619
```

## Hypothesis Testing: the T-Statistic

Reproduce the example from the presentation

```
# define parameters
mu <- 39000
sample_mean <- 37000
sample_sd <- 6150
n <- 100

# calculate the standard error of the sample mean
se <- sample_sd / sqrt(n)

# compute the t-statistic (and compare it with the critical value)
t <- (sample_mean - mu) / se

# look up the p-value
# (the fraction of the mass under the standard normal distribution)
2*pnorm(-abs(t))

## [1] 0.001145829
```

### Extended Example

```
# I) Compute the t-value step by step with our own implementation while controlling the properties of t

# define size of sample
n <- 100
# draw the random sample from a normal distribution with mean 10 and sd 2
sample <- rnorm(n, mean = 10, sd = 2)

# compute the sample mean
sample_mean <- mean(sample)
# compute the sample sd
sample_sd <- sd(sample)
# estimated standard error of the mean
mean_se <- sample_sd/sqrt(length(sample))

# compute the t-statistic for the null hypothesis: H0: mu = 9
t <- (sample_mean - 10) / mean_se
t

## [1] 0.9884692

# get the p value
2*pnorm(-abs(t))

## [1] 0.3229229

# II) Apply the R-function t.test to do the same!
t.test(sample, mu = 10)

##
## One Sample t-test
```

```
##
## data:  sample
## t = 0.98847, df = 99, p-value = 0.3253
## alternative hypothesis: true mean is not equal to 10
## 95 percent confidence interval:
##  9.80395 10.58528
## sample estimates:
## mean of x
## 10.19462
```

## Exercises

1. Got to <http://stat.ethz.ch/R-manual/R-devel/library/datasets/html/00Index.html>. 2. Pick a data set that interests you, load it with the `data()` function.

```
# yes, I choose the swiss data set...
data(swiss)
# have a look at what we are dealing with here
str(swiss)

## 'data.frame':  47 obs. of  6 variables:
## $ Fertility      : num  80.2 83.1 92.5 85.8 76.9 76.1 83.8 92.4 82.4 82.9 ...
## $ Agriculture    : num  17 45.1 39.7 36.5 43.5 35.3 70.2 67.8 53.3 45.2 ...
## $ Examination    : int  15 6 5 12 17 9 16 14 12 16 ...
## $ Education      : int  12 9 5 7 15 7 7 8 7 13 ...
## $ Catholic       : num  9.96 84.84 93.4 33.77 5.16 ...
## $ Infant.Mortality: num  22.2 22.2 20.2 20.3 20.6 26.6 23.6 24.9 21 24.4 ...

# have a look at the first few lines
head(swiss)

##           Fertility Agriculture Examination Education Catholic
## Courtelary      80.2         17.0           15          12      9.96
## Delemont        83.1         45.1            6           9     84.84
## Franches-Mnt    92.5         39.7            5           5     93.40
## Moutier         85.8         36.5           12           7     33.77
## Neuveville      76.9         43.5           17          15      5.16
## Porrentruy      76.1         35.3            9           7     90.57
##           Infant.Mortality
## Courtelary             22.2
## Delemont               22.2
## Franches-Mnt           20.2
## Moutier                20.3
## Neuveville             20.6
## Porrentruy             26.6

# get more info about the variables in that data set
?swiss
```

3. Pose an empirical question you want to answer with that data set. In the case of the `swiss` data set, for example, is the average fertility in the entire population 85?
4. Formulate a null-hypothesis and test it with the techniques learned today.  $H_0: \mu = 85$ . This requires a t-statistic of the Fertility mean. We can compute this statistic either with the formulas and R-functions discussed on Day 2. Or we directly compute it with `t.test()` provided by R. The following code demonstrates both approaches step by step.

## I. Own implementation (with a detailed explanation)

```
# a) compute sample mean and sample standard deviation, record how many observations we are having
# in our sample, define the population mean (that you want to test for)
sample_mean <- mean(swiss$Fertility)
sample_sd <- sd(swiss$Fertility)
n <- length(swiss$Fertility) # alternatively use nrow(swiss)
mu <- 85

# b) compute the (estimate of the) sample mean standard error
se <- sample_sd / sqrt(n)

# c) compute the t-statistic
t <- (sample_mean - mu) / se
t
```

```
## [1] -8.154018
```

```
# d) check what p-value is associated with that t-statistic
# i.e., check what fraction of the standard normal distribution has an at least as extreme value as
# the t value we computed.
pval <- 2*pnorm(-abs(t))
pval
```

```
## [1] 3.520284e-16
```

Note how we get from the t-value to the p-value here: `2*pnorm(-abs(t))`. This short line of code involves a lot. Let's have a look at specific aspects. Recall the standard normal distribution function which is implemented in `pnorm(x)`. This function returns the fraction of the mass under the standard normal distribution curve that is smaller than or equal to `x` (i.e., the probability of observing a value that is equal to or smaller than `x`). Let's have a closer look in R:

```
# the probability of observing a value at least as small as -1 in a standard normal distribution
pnorm(-1)
```

```
## [1] 0.1586553
```

```
# or in percent
pnorm(-1) * 100
```

```
## [1] 15.86553
```

Recall that we are interested in the probability of observing a value at least as *extreme* as `t`. Thus we need to consider the absolute value of `t`, `|t|`, and want to know at once what the probability of observing a value at least as small as `-t` or at least as big as `t` (therefore, the `abs(t)` which returns the absolute value).

```
# demonstration of concept of absolute value
abs(-1)
```

```
## [1] 1
```

```
abs(1)
```

```
## [1] 1
```

```
# in our example of t
abs(t)
```

```
## [1] 8.154018
```

Recall that `pnorm(x)` returns by default the probability of observing a value that is at least as small as `x` therefore the `-abs(t)`. Finally, we also know that the standard normal distribution is symmetric. Thus the fraction of mass under the curve between `-abs(t)` and `-infinity` (everything to the left of `-abs(t)` under the curve) is exactly the same as the one between `abs(t)` and `infinity` (everything to the left of `abs(t)` under the curve). Hence we can simply multiply the value of `pnorm(-abs(t))` by two: `2*pnorm(-abs(t))` for simplicity's sake. Alternatively we could compute the probability for the 'lower tail' and the 'upper tail' separately and build the sum of the two.

```
# compute the lower tail probability
lowerp <- pnorm(-abs(t))
# compute the upper tail probability
upperp <- pnorm(abs(t), lower.tail = FALSE)
# build the two-tail p-value
twotailp <- lowerp + upperp

# proof that this is the same as above
twotailp == 2*pnorm(-abs(t))
```

```
## [1] TRUE
```

## II. Apply the R function for t-tests (i.e., the `tl;dr` version...)

```
# t-test for H0: mu = 85
t.test(swiss$Fertility, mu = 85)

##
## One Sample t-test
##
## data:  swiss$Fertility
## t = -8.154, df = 46, p-value = 1.755e-10
## alternative hypothesis: true mean is not equal to 85
## 95 percent confidence interval:
##  66.47485 73.81025
## sample estimates:
## mean of x
##  70.14255
```

Given this result, we can quite confidently reject  $H_0$  and conclude that the population mean of this fertility measure is very unlikely 85.